

Next-Gen AI for Next-Gen Gaming

GLM-5.1: #1 on SWE-Bench Pro | 744B MoE | 3.5× cheaper than Claude Opus | Apache 2.0 Open Source

#1

SWE-Bench Pro
Globally

94.6%

of Claude Opus 4.6
Coding Performance

8h

Autonomous
Task Duration

3.5×

Lower Cost
vs. Claude Opus

● AVAILABLE NOW

GLM-5.1

Flagship · 744B MoE · 200K context

- ◆ #1 SWE-Bench Pro — real-world GitHub debugging
- ◆ 94.6% of Claude Opus 4.6 coding performance
- ◆ 8-hour autonomous: plan → execute → deliver
- ◆ 58.4% SWE-Bench Pro, beats GPT-5.4 & Opus 4.6
- ◆ Open-source Apache 2.0 — deploy anywhere
- ◆ 28% coding improvement over GLM-5 baseline

● COMING SOON

GLM New Version

Next-Gen · 1M Context · Enhanced Reasoning

- ◆ Further coding gains on #1 SWE-Bench Pro base
- ◆ Enhanced math, science & logic reasoning
- ◆ Advanced multimodal: image, video, 3D
- ◆ Longer autonomous runs beyond 8 hours
- ◆ Game-specific fine-tuning pipeline
- ◆ API-compatible with GLM-5.1 endpoints

BENCHMARK COMPARISON

Benchmark	GLM-5.1	Claude Opus 4.6	GPT-5.4	DeepSeek V3.2
SWE-Bench Pro	58.4%	53.2%	52.1%	38.7%
SWE-Bench Verified	77.8%	75.6%	72.8%	62.3%
Terminal-Bench 2.0	60.7%	58.9%	56.4%	45.2%
Text Arena (Overall)	#4 Global	#1 Global	#2 Global	#9 Global
Open Source	✓ Apache 2.0	✗	✗	✓ MIT

GAMING USE CASES FOR RIOT

■ NPC Dialogue & Lore

200K context handles entire champion rosters and world-building in one pass

■ Content Pipeline

Automate scripts, patch notes, social.
GLM-5V generates visuals from text.

■ Anti-Cheat & Toxicity

Detect toxic language across 50+ langs
with high precision and low latency.

■ Game Dev Automation

#1 SWE-Bench Pro: 8h autonomous code debug, refactor, entire feature branches.

■ Player Analytics

NL query of telemetry: instant SQL gen.
Win rate anomalies, churn prediction.

■ Localization at Scale

Context-aware translation for 20+ langs.
Gaming jargon, champion names, abilities.

GLM-5.1 Pricing (Z.AI)

Input	\$1.40 /1M tokens
Cached Input	\$0.26 /1M tokens (80% off)
Output	\$4.40 /1M tokens
Context	200K tokens
GLM New Version pricing: TBD (competitive) · 1M context	

Cost Comparison @ Scale (10B tokens/mo)

	GLM-5.1	Claude Opus 4.6	Savings
Input cost	\$14,000	\$150,000	91%
Output cost	\$44,000	\$750,000	94%
Monthly total	\$58,000	\$900,000	\$842K/mo